

▾ **Actividad en Clase Transfer Learning**

Victor Hugo Arreola Elenes - A01635682

Fernando Ojeda Marín - A01639252

Mario Alberto Castañeda Martínez - A01640152

```
1 import matplotlib.pyplot as plt
2 import numpy as np
3 import os
4 import tensorflow as tf

1 from google.colab import drive
2 drive.mount('/content/drive')

    Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_remount=True).

1 import os
2
3 # Ruta a la carpeta en Google Drive donde se encuentran tus datos
4 PATH = '/content/drive/MyDrive/Tec/7mo Semestre/Inteligencia Artificial II/Deep Learning/Data-Monitor-Teclados-Mouse/'
5
6 # Rutas a las carpetas de entrenamiento y validación
7 train_dir = os.path.join(PATH, 'train')
8 validation_dir = os.path.join(PATH, 'validation')
9
10 BATCH_SIZE = 32
11 IMG_SIZE = (128,128)
12
13 train_dataset = tf.keras.utils.image_dataset_from_directory(train_dir,
14                                                             shuffle = True,
15                                                             batch_size=BATCH_SIZE,
16                                                             image_size=IMG_SIZE)
17
18 validation_dataset = tf.keras.utils.image_dataset_from_directory(validation_dir,
19                                                                    shuffle = True,
20                                                                    batch_size=BATCH_SIZE,
21                                                                    image_size=IMG_SIZE)

    Found 732 files belonging to 3 classes.
    Found 93 files belonging to 3 classes.

1 class_names = train_dataset.class_names
2
3 plt.figure(figsize=(10,10))
4 for image, labels in train_dataset.take(1):
5     for i in range(9):
6         ax=plt.subplot(3,3,i+1)
7         plt.imshow(image[i].numpy().astype('uint8'))
8         plt.title(class_names[labels[i]])
9         plt.axis("off")
```

```
mouse keyboard keyboard

1 val_batches = tf.data.experimental.cardinality(validation_dataset)
2 test_dataset = validation_dataset.take(val_batches // 5)
3 validation_dataset = validation_dataset.skip(val_batches // 5)
4
5 print("numero de batches para test_data set = %d" %tf.data.experimental.cardinality(test_dataset))
6 print("numero de batches para validation_data set = %d" %tf.data.experimental.cardinality(validation_dataset))
```

numero de batches para test_data set = 0
numero de batches para validation_data set = 3

```
AUTOTUNE = tf.data.AUTOTUNE
2
3 train_dataset = train_dataset.prefetch(buffer_size=AUTOTUNE)
4 validation_dataset = validation_dataset.prefetch(buffer_size=AUTOTUNE)
5 test_dataset = test_dataset.prefetch(buffer_size=AUTOTUNE)
```

```
1 data_augmentation = tf.keras.Sequential([
2     tf.keras.layers.RandomFlip("horizontal"),
3     tf.keras.layers.RandomRotation(0.2),
4 ])
```

```
1 for image, _ in train_dataset.take(1):
2     plt.figure(figsize=(10,10))
3     first_image = image[0]
4     for image, labels in train_dataset.take(1):
5         for i in range(9):
6             ax=plt.subplot(3,3,i+1)
7             augmented_image = data_augmentation(tf.expand_dims(first_image,0))
8             plt.imshow(augmented_image[0]/255)
9             plt.axis("off")
```



```
1 rescale = tf.keras.layers.Rescaling(1./127.5,offset = 1)
2
3 preprocess_input = tf.keras.applications.mobilenet_v2.preprocess_input

1 IMG_SHAPE = IMG_SIZE + (3,)
2 print(IMG_SHAPE)
3
4 base_model = tf.keras.applications.InceptionV3(input_shape=IMG_SHAPE,include_top=False,weights="imagenet")

(128, 128, 3)

1 image_batch, label_batch = next(iter(train_dataset))
2 feature_batch = base_model(image_batch)
3
4 print(feature_batch.shape)

(32, 2, 2, 2048)
```

```
1 base_model.trainable=False
2 base_model.summary()
    batch_normalization_454 (Batch Normalization)      1152      [conv2d_454[0][0]]
    batch_normalization_455 (Batch Normalization)      1152      ['conv2d_455[0][0]']
    batch_normalization_458 (Batch Normalization)      1152      ['conv2d_458[0][0]']
    batch_normalization_459 (Batch Normalization)      1152      ['conv2d_459[0][0]']
    conv2d_460 (Conv2D)                                245760     ['average_pooling2d_43[0][0]']
    batch_normalization_452 (Batch Normalization)       960        ['conv2d_452[0][0]']
    activation_454 (Activation)                         0          ['batch_normalization_454[0][0]']
    activation_455 (Activation)                         0          ['batch_normalization_455[0][0]']
    activation_458 (Activation)                         0          ['batch_normalization_458[0][0]']
    activation_459 (Activation)                         0          ['batch_normalization_459[0][0]']
    batch_normalization_460 (Batch Normalization)       576        ['conv2d_460[0][0]']
    activation_452 (Activation)                         0          ['batch_normalization_452[0][0]']
    mixed9_0 (Concatenate)                             0          ['activation_454[0][0]',
    'activation_455[0][0]']
    concatenate_8 (Concatenate)                       0          ['activation_458[0][0]',
    'activation_459[0][0]']
    activation_460 (Activation)                         0          ['batch_normalization_460[0][0]']
    mixed9 (Concatenate)                              0          ['activation_452[0][0]',
    'mixed9_0[0][0]',
    'concatenate_8[0][0]',
    'activation_460[0][0]']
    conv2d_465 (Conv2D)                                917504     ['mixed9[0][0]']
    batch_normalization_465 (Batch Normalization)      1344       ['conv2d_465[0][0]']
    activation_465 (Activation)                         0          ['batch_normalization_465[0][0]']
    conv2d_462 (Conv2D)                                786432     ['mixed9[0][0]']
```

```
1 global_average_layer = tf.keras.layers.GlobalAveragePooling2D()
2 feature_batch_average = global_average_layer(feature_batch)
3 print(feature_batch_average.shape)

(32, 2048)

1 prediction_layer = tf.keras.layers.Dense(3)
2 prediction_batch = prediction_layer(feature_batch_average)
3 print(prediction_batch.shape)

(32, 3)
```

Unir Modelo

```
1 from tensorflow.keras.callbacks import ModelCheckpoint

1 inputs = tf.keras.Input(shape=(128, 128, 3))
2 x = data_augmentation(inputs)
3 x = preprocess_input(x)
4 x = base_model(x, training=False)
5 x = global_average_layer(x)
6 x = tf.keras.layers.Dropout(0.2)(x)
7
8 outputs = prediction_layer(x)
9
10 model = tf.keras.Model(inputs, outputs)
11

1 base_learning_rate = 0.00001
2 model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=base_learning_rate),
3               loss=tf.keras.losses.SparseCategoricalCrossentropy(),
4               metrics=['accuracy'])
5 model.summary()
```

Model: "model_4"

Layer (type)	Output Shape	Param #
input_10 (InputLayer)	[(None, 128, 128, 3)]	0

```
sequential_4 (Sequential)   (None, 128, 128, 3)      0

tf.math.truediv_4 (TFOpLam  (None, 128, 128, 3)      0
bda)

tf.math.subtract_4 (TFOpLa  (None, 128, 128, 3)      0
mbda)

inception_v3 (Functional)   (None, 2, 2, 2048)       21802784

global_average_pooling2d_4  (None, 2048)             0
(GlobalAveragePooling2D)

dropout_4 (Dropout)         (None, 2048)             0

dense_5 (Dense)             (None, 3)                6147

=====
Total params: 21808931 (83.19 MB)
Trainable params: 6147 (24.01 KB)
Non-trainable params: 21802784 (83.17 MB)
=====

1 initial_epochs = 20
2 loss0, accuracy0 = model.evaluate(validation_dataset)
3 print(loss0)
4 print(accuracy0)

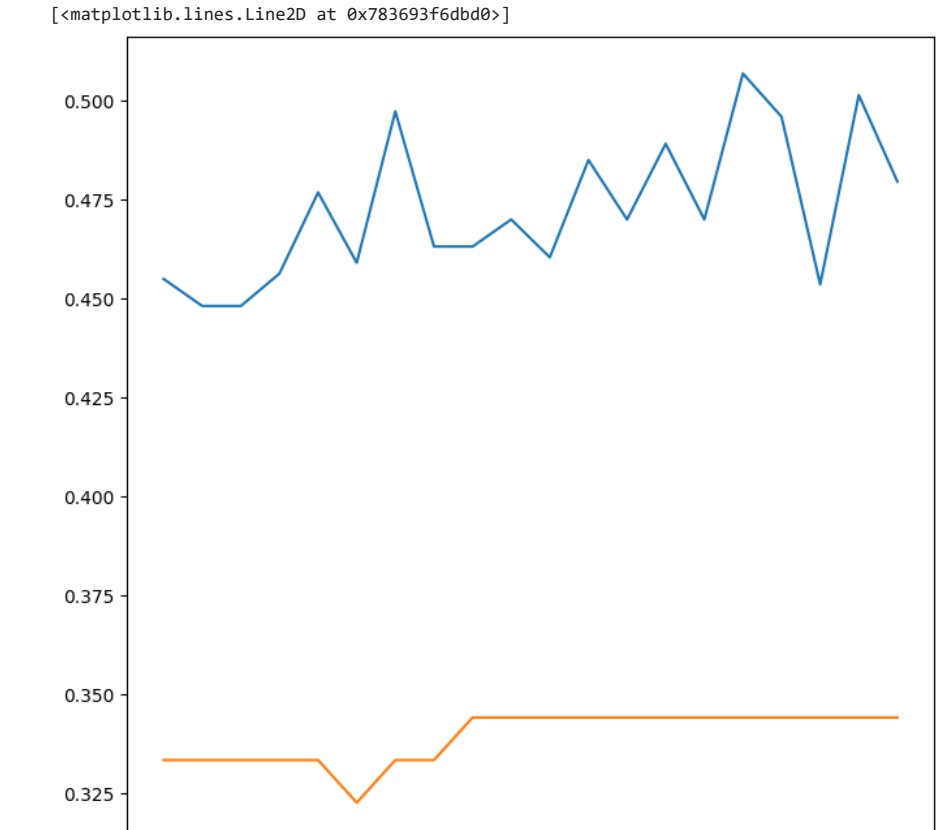
3/3 [=====] - 6s 58ms/step - loss: 6.5696 - accuracy: 0.3333
6.569586753845215
0.3333333432674408

1 history = model.fit(train_dataset,
2                     epochs = initial_epochs,
3                     validation_data = validation_dataset)

Epoch 9/20
23/23 [=====] - ETA: 0s - loss: 5.4831 - accuracy: 0.4399
Epoch 9: val_accuracy did not improve from 0.32258
23/23 [=====] - 2s 73ms/step - loss: 5.4831 - accuracy: 0.4399 - val_loss: 6.6910 - val_accuracy: 0.3226
Epoch 10/20
23/23 [=====] - ETA: 0s - loss: 5.4192 - accuracy: 0.4385
Epoch 10: val_accuracy did not improve from 0.32258
23/23 [=====] - 2s 80ms/step - loss: 5.4192 - accuracy: 0.4385 - val_loss: 6.6904 - val_accuracy: 0.3226
Epoch 11/20
23/23 [=====] - ETA: 0s - loss: 5.4576 - accuracy: 0.4112
Epoch 11: val_accuracy did not improve from 0.32258
23/23 [=====] - 2s 58ms/step - loss: 5.4576 - accuracy: 0.4112 - val_loss: 6.6890 - val_accuracy: 0.3226
Epoch 12/20
22/23 [=====>..] - ETA: 0s - loss: 5.5100 - accuracy: 0.4403
Epoch 12: val_accuracy did not improve from 0.32258
23/23 [=====] - 2s 57ms/step - loss: 5.5302 - accuracy: 0.4440 - val_loss: 6.6991 - val_accuracy: 0.3226
Epoch 13/20
22/23 [=====>..] - ETA: 0s - loss: 5.5067 - accuracy: 0.4190
Epoch 13: val_accuracy did not improve from 0.32258
23/23 [=====] - 2s 62ms/step - loss: 5.5250 - accuracy: 0.4153 - val_loss: 6.6972 - val_accuracy: 0.3226
Epoch 14/20
22/23 [=====>..] - ETA: 0s - loss: 5.0702 - accuracy: 0.4673
Epoch 14: val_accuracy did not improve from 0.32258
23/23 [=====] - 2s 56ms/step - loss: 5.1084 - accuracy: 0.4631 - val_loss: 6.6960 - val_accuracy: 0.3226
Epoch 15/20
22/23 [=====>..] - ETA: 0s - loss: 5.3358 - accuracy: 0.4105
Epoch 15: val_accuracy did not improve from 0.32258
23/23 [=====] - 2s 57ms/step - loss: 5.3875 - accuracy: 0.4098 - val_loss: 6.6949 - val_accuracy: 0.3226
Epoch 16/20
22/23 [=====>..] - ETA: 0s - loss: 5.4233 - accuracy: 0.4304
Epoch 16: val_accuracy did not improve from 0.32258
23/23 [=====] - 2s 75ms/step - loss: 5.4437 - accuracy: 0.4331 - val_loss: 6.6944 - val_accuracy: 0.3226
Epoch 17/20
23/23 [=====] - ETA: 0s - loss: 4.8749 - accuracy: 0.4467
Epoch 17: val_accuracy did not improve from 0.32258
23/23 [=====] - 2s 71ms/step - loss: 4.8749 - accuracy: 0.4467 - val_loss: 6.6938 - val_accuracy: 0.3226
Epoch 18/20
23/23 [=====] - ETA: 0s - loss: 4.9553 - accuracy: 0.4522
Epoch 18: val_accuracy did not improve from 0.32258
23/23 [=====] - 2s 61ms/step - loss: 4.9553 - accuracy: 0.4522 - val_loss: 6.6859 - val_accuracy: 0.3226
Epoch 19/20
23/23 [=====] - ETA: 0s - loss: 5.0189 - accuracy: 0.4194
Epoch 19: val_accuracy did not improve from 0.32258
23/23 [=====] - 2s 58ms/step - loss: 5.0189 - accuracy: 0.4194 - val_loss: 6.6712 - val_accuracy: 0.3226
Epoch 20/20
22/23 [=====>..] - ETA: 0s - loss: 4.8692 - accuracy: 0.4489
Epoch 20: val_accuracy did not improve from 0.32258
23/23 [=====] - 2s 57ms/step - loss: 4.8606 - accuracy: 0.4508 - val_loss: 6.5687 - val_accuracy: 0.3226
Epoch 1/20
23/23 [=====] - 29s 143ms/step - loss: 5.1122 - accuracy: 0.4549 - val_loss: 6.5584 - val_accuracy: 0.3333
Epoch 2/20
23/23 [=====] - 3s 113ms/step - loss: 5.1438 - accuracy: 0.4481 - val_loss: 6.5413 - val_accuracy: 0.3333
Epoch 3/20
23/23 [=====] - 2s 79ms/step - loss: 5.1721 - accuracy: 0.4481 - val_loss: 6.4413 - val_accuracy: 0.3333
Epoch 4/20
23/23 [=====] - 2s 83ms/step - loss: 5.0210 - accuracy: 0.4563 - val_loss: 6.4584 - val_accuracy: 0.3333
Epoch 5/20
23/23 [=====] - 3s 95ms/step - loss: 4.6242 - accuracy: 0.4768 - val loss: 6.4706 - val accuracy: 0.3333

1 acc = history.history['accuracy']
2 val_acc = history.history['val_accuracy']
3
4 plt.figure(figsize=(8,8))
5 plt.plot(acc, label = 'Training acc')
6 plt.plot(val_acc, label = 'Validation acc')
```





Fine Tuning

```
1 base_model.trainable = True
2
3 print("Numero de capas ", len(base_model.layers))

Numero de capas  311

1 fine_tune_at = 100
2 for layer in base_model.layers[:fine_tune_at]:
3     layer.trainable = False

1 model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=base_learning_rate/10),
2               loss=tf.keras.losses.SparseCategoricalCrossentropy(),
3               metrics=['accuracy'])
4
5 model.summary()
```

Model: "model_4"

Layer (type)	Output Shape	Param #
input_10 (InputLayer)	[(None, 128, 128, 3)]	0
sequential_4 (Sequential)	(None, 128, 128, 3)	0
tf.math.truediv_4 (TFOpLambda)	(None, 128, 128, 3)	0
tf.math.subtract_4 (TFOpLambda)	(None, 128, 128, 3)	0
inception_v3 (Functional)	(None, 2, 2, 2048)	21802784
global_average_pooling2d_4 (GlobalAveragePooling2D)	(None, 2048)	0
dropout_4 (Dropout)	(None, 2048)	0
dense_5 (Dense)	(None, 3)	6147

=====

Total params: 21808931 (83.19 MB)

Trainable params: 19632515 (74.89 MB)

Non-trainable params: 2176416 (8.30 MB)

=====

```
1 fine_tune_epochs = 10
2 total_epochs = initial_epochs + fine_tune_epochs
3
4 history_fine = model.fit(train_dataset,
5                           epochs = total_epochs,
6                           initial_epoch = history.epoch[-1],
7                           validation_data = validation_dataset)
```

Epoch 20/30

23/23 [=====] - 33s 153ms/step - loss: 4.5181 - accuracy: 0.4891 - val_loss: 6.8590 - val_accuracy: 0.3548

Epoch 21/30

23/23 [=====] - 2s 78ms/step - loss: 4.5123 - accuracy: 0.4645 - val_loss: 6.9533 - val_accuracy: 0.3548

Epoch 22/30

23/23 [=====] - 2s 84ms/step - loss: 4.4939 - accuracy: 0.4891 - val_loss: 6.8377 - val_accuracy: 0.3656

Epoch 23/30

23/23 [=====] - 2s 80ms/step - loss: 4.8037 - accuracy: 0.4658 - val_loss: 6.8326 - val_accuracy: 0.3763

Epoch 24/30

23/23 [=====] - 3s 105ms/step - loss: 4.5162 - accuracy: 0.4959 - val_loss: 6.8186 - val_accuracy: 0.3763

Epoch 25/30

23/23 [=====] - 3s 88ms/step - loss: 4.3260 - accuracy: 0.4918 - val_loss: 6.8142 - val_accuracy: 0.3763

```
Epoch 26/30
23/23 [=====] - 2s 79ms/step - loss: 4.4052 - accuracy: 0.5109 - val_loss: 6.8108 - val_accuracy: 0.3763
Epoch 27/30
23/23 [=====] - 2s 78ms/step - loss: 3.7747 - accuracy: 0.4986 - val_loss: 6.7981 - val_accuracy: 0.3763
Epoch 28/30
23/23 [=====] - 2s 78ms/step - loss: 4.1048 - accuracy: 0.5191 - val_loss: 6.7978 - val_accuracy: 0.3763
Epoch 29/30
23/23 [=====] - 2s 85ms/step - loss: 3.9071 - accuracy: 0.5301 - val_loss: 6.7940 - val_accuracy: 0.3763
Epoch 30/30
23/23 [=====] - 3s 104ms/step - loss: 4.1700 - accuracy: 0.5000 - val_loss: 6.7941 - val_accuracy: 0.3763
```

```
1 from tensorflow import keras
2 model = keras.models.load_model('mejor_modelo.h5')
```

```
1 n=10 ##Number of image
2 plt.figure(figsize=(2,2))
3 plt.imshow(image[n].numpy().astype('uint8'))
4 plt.title(class_names[labels[n]])
5 plt.axis("off")
```

(-0.5, 127.5, 127.5, -0.5)

keyboard



```
1 import numpy as np
2 predictions = model.predict(image)
3 print(predictions[n])
4 print("This image most likely belongs to {} with a {:.2f} percent confidence.".format(class_names[np.argmax(predictions[n])], 10 * np.max(
5
```

```
1/1 [=====] - 1s 1s/step
[ 1.1176975 -1.4487779 -1.0702685]
This image most likely belongs to keyboard with a 11.18 percent confidence.
```

```
1 n=3 ##Number of image
2 plt.figure(figsize=(2,2))
3 plt.imshow(image[n].numpy().astype('uint8'))
4 plt.title(class_names[labels[n]])
5 plt.axis("off")
```

(-0.5, 127.5, 127.5, -0.5)

mouse



```
1 import numpy as np
2 predictions = model.predict(image)
3 print(predictions[n])
4 print("This image most likely belongs to {} with a {:.2f} percent confidence.".format(class_names[np.argmax(predictions[n])], 10 * np.max(
5
```

```
1/1 [=====] - 0s 43ms/step
[ 1.6942713 -0.07480229 2.0224931 ]
This image most likely belongs to mouse with a 20.22 percent confidence.
```

```
1 n=20 ##Number of image
2 plt.figure(figsize=(2,2))
3 plt.imshow(image[n].numpy().astype('uint8'))
4 plt.title(class_names[labels[n]])
5 plt.axis("off")
```

(-0.5, 127.5, 127.5, -0.5)

monitor



```
1 import numpy as np
2 predictions = model.predict(image)
3 print(predictions[n])
4 print("This image most likely belongs to {} with a {:.2f} percent confidence.".format(class_names[np.argmax(predictions[n])], 10 * np.max(
5
```

```
1/1 [=====] - 0s 29ms/step
[ 0.00783599 0.5664257 -0.07600397]
```

This image most likely belongs to monitor with a 5.66 percent confidence.